

Python 图像处理基础



胡俊峰 2020/04/20

北京大学信息科学技术学院



内容

- 安装CV软件包
- Python的图像格式与存储方案
- 简单特征统计与直方图变换
- 卷积、角点
- 图像特征提取

```
3 import cv2 # yes, we are using opencv version 3
```

```
: 1 # code to find version of opencv  
2 cv2.__version__
```

Conda install opencv

```
'3.4.2'
```

```
: 1 img_dir = 'common/'  
2 messi_gray = cv2.imread(img_dir+'data/messi.jpg', 0) # 第二个参数 0 grey  
3 my_gshow(plt.gca(), messi_gray) # 图片-线... 都是 axes 的子对象  
4 messi_gray
```


```
array([[ 43,  46,  48, ...,  55,  53,  50],  
       [ 41,  46,  50, ...,  60,  58,  55],  
       [ 46,  51,  56, ...,  64,  63,  60],  
       ...,  
       [120, 110, 107, ..., 113, 114, 124],  
       [116, 119, 108, ..., 111, 122, 117],  
       [107, 118, 129, ..., 104, 105, 104]], dtype=uint8)
```



➤ Matplotlib get current axis


```
1 xs = np.linspace(0,10, 100) # 100 evenly spaced points from 0 to 10 inclusive
2 ys = np.exp(xs) #  $y = e^x$ 
3
4 # fig = plt.gcf() # explicitly get the default (current) figure (holder for "full" graphic) (rarely needed)
5 ax = plt.gca() # explicitly get the default (current) drawing axis
6
7 ax.plot(xs, ys) # basic graph from two sequences: x-points and y-points
8 ax.set_title("Example Graph")
9 ax.set_xlabel("Time")
10 # ax.axis('off'); # uncomment me to see the difference with/without spines and labels
```

Text(0.5, 0, 'Time')



```
1 ax = plt.gca()
2 arr = npr.randint(0,10,20)
3 ax.hist(arr, bins=10)
4
5 ax.set_xlabel("Integer")
6 ax.set_ylabel("Count")
7 ax.set_title("Histogram of Integers")
8
9 import collections as co
10 print(co.Counter(arr)) # pure python counts of occurrences
```

Counter({3: 5, 0: 3, 6: 3, 8: 2, 9: 2, 1: 1, 2: 1, 4: 1, 5: 1, 7: 1})



```
1 # note: by default the range [0,1] is expanded to fill the [0,255] intensity scale
2 #       so: 0--->0 (black) and 1--->255 (white)
3
4 # also: remove the x/y grid points (matplotlib calls these "spines") and the frame
5 ax = plt.gca()
6 ax.imshow(arr, cmap='gray')
7 ax.axis('off');
```



定义两个显示用函数




```
1  # line 5:  def my_show(**kwargs) --> takes any "extra" keyword arguments and
2  #                                                  puts them in a dictionary named kwargs
3  # line 7:  ax.imshow(**kwargs)    --> takes kwargs (a dictionary) and
4  #                                                  "expands" them into keyword arguments to imshow
5  def my_show(ax, img, title=None, interpolation='bicubic', **kwargs): ← 采用二次插值填充
6      ' helper to display an image on an axes without grid/spine '
7      ax.imshow(img, interpolation = interpolation, **kwargs)
8      ax.axis('on')
9      if title:
10         ax.set_title(title)
11
12 def my_gshow(ax, img, title=None, cmap='gray', interpolation='bicubic', **kwargs):
13     ' helper to display an image, in grayscale, on an axes without grid/spine '
14     my_show(ax, img, title=title, cmap='gray', interpolation=interpolation, **kwargs)
```

```

1 fig, axes = plt.subplots(1,3,figsize=(9,3)) # 1 row, 3 columns of figures
2 print(fig, axes)
3 # axes is a numpy array with nrows, ncols ... unless one of the values is 1 ... then it is a
4 # 1-D thing (but see squeeze=True squeeze=False argument to subplots to keep dimensions)
5 arr = np.array([[0,1],
6                 [1,0]], dtype=np.float64)
7 my_show(axes[0], arr)
8 my_gshow(axes[1], arr) ← 一组axis
9 my_gshow(axes[2], arr, interpolation=None)

```

Figure(648x216) [<matplotlib.axes._subplots.AxesSubplot object at 0x000001E74CE145F8>
 <matplotlib.axes._subplots.AxesSubplot object at 0x000001E74CE522B0>
 <matplotlib.axes._subplots.AxesSubplot object at 0x000001E74CE79940>]



加载显示图像

```
1 img_dir = 'common/'
2 messi_gray = cv2.imread(img_dir+'data/messi.jpg', 0) # 第二个参数 0 grey
3 my_gshow(plt.gca(), messi_gray) # 参数1: axis 参数2: array
4 messi_gray
```

```
array([[ 43,  46,  48, ...,  55,  53,  50],
       [ 41,  46,  50, ...,  60,  58,  55],
       [ 46,  51,  56, ...,  64,  63,  60],
       ...,
       [120, 110, 107, ..., 113, 114, 124],
       [116, 119, 108, ..., 111, 122, 117],
       [107, 118, 129, ..., 104, 105, 104]], dtype=uint8)
```



```

1 messi_color = cv2.imread(img_dir+'data/messi.jpg') # default flag is 1 "color"
2 print(type(messi_color), messi_color.shape, messi_color.dtype)
3 my_show(plt.gca(), messi_color)
4 # 实际编码 GBR - 习惯 RGB ← messi_rgb = cv2.cvtColor(messi_color, cv2.COLOR_BGR2RGB)

```

<class 'numpy.ndarray'> (342, 548, 3) uint8



```

1 # opencv is GBR; matplotlib is RGB.
2 my_show(plt.gca(), messi_color[:, :, ::-1]) # walk last axis in opposite order (we'll never do this again!)

```



```
1 # we can also use scipy to read in images:
2 from scipy import ndimage
3 img = ndimage.imread(img_dir+'data/messi.jpg') # default is 1 RGB
4 #img1 = plt.imread(img_dir+'data/messi.jpg')
5 my_show(plt.gca(), img1)
6 img1.shape
```

C:\Users\hjf_p\Anaconda3\lib\site-packages\ipykernel_launcher.py:3: DeprecationWarning: `imread` is deprecated!
`imread` is deprecated in SciPy 1.0.0.
Use ``matplotlib.pyplot.imread`` instead.
This is separate from the ipykernel package so we can avoid doing imports until

(342, 548, 3)




```
1  # these come up frequently. we'll always want rgb (instead of bgr)
2  # and we often need both rgb and grayscale (grayscale starts many processing steps)
3  def my_read(filename):
4      ' read from an image file to an rgb '
5      img = cv2.imread(filename)
6      return cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
7
8  def my_read_cg(filename):
9      ' read from an image file to an rgb and a grayscale image array '
10     rgb = my_read(filename)
11     gray = cv2.cvtColor(rgb, cv2.COLOR_RGB2GRAY)
12     return rgb, gray ← 返回一个tuple
13
14 # now we can do this:
15 messi_rgb = my_read(img_dir+'data/messi.jpg')
16
17 # or if we need both
18 messi_rgb, messi_gray = my_read_cg(img_dir+'data/messi.jpg') ←
```

```
1 messi_rgb = my_read(img_dir+'data/messi.jpg') # 读入时直接转换RGB
```

Since messi_rgb is "just" a NumPy array, we can do NumPy array things:


```
1 print(messi_rgb[100, 100],      # access a pixel
2       messi_rgb[300, :, :].shape) # sub-select a row; it's an array also.  take
```

[200 166 156] (548, 3) ← 一个三通道像素，每行的数据为548*3的矩阵

```
1 # pixels are people ... err ... arrays too
2 pixel = messi_rgb[100, 100]
3 print(type(pixel),
4       pixel.shape, # 1-D, scalar, array
5       pixel)
```

<class 'numpy.ndarray'> (3,) [200 166 156]

```
1 # messi's right wrist has a white spot!  
2 messi_rgb[100:105, 100:105]=[255, 255, 255] # white (note, our target pixel also had 3 spots  
3 my_show(plt.gca(), messi_rgb)
```



```
1 ball_soi = messi_rgb[280:340, 330:390] # "soi" = square of interest :)
2 messi_rgb[273:333, 100:160] = ball_soi # copy to new area
3 my_show(plt.gca(), messi_rgb)
```




```
1 # often we want to access color channels separately
2 # split to separate arrays per color (costly, prefer to access by indexing)
3 chans = r, g, b = cv2.split(messi_rgb)
4 restored = cv2.merge((r, g, b))
```

```
1 fig, axes = plt.subplots(1, 4, figsize=(12, 3))
2 axes = axes.flat # a numpy array of axes
3
4 # handle first as special case
5 first_axis = next(axes)
6 my_show(first_axis, messi_rgb)
7 first_axis.set_title("original")
8
9 # display per channel images
10 for ax, ch, name in zip(axes, chans, ["R", "G", "B"]):
11     my_gshow(ax, g)
12     ax.set_title("{} channel".format(name))
```

original



R channel



G channel




B channel




```

1  # we can often just use indexing directly (see line 9)
2  # also show off matplotlib histograms
3
4  color_to_index = {"R":0, "G":1, "B":2} # map strings to appropriate index in
5
6  fig, axes = plt.subplots(1,3,figsize=(12,3), sharey=True)
7  for ax, color in zip(axes, color_to_index):
8      c = color_to_index[color]
9      this_channel = messi_rgb[:, :, c].ravel() # 1D flat array view without copying
10
11     ax.hist(this_channel, 256, normed=True)
12     ax.set_title("Histogram for {}".format(color))

```




```
1 ml = my_read(img_dir+' /data/ml.png')
2 frog = my_read(img_dir+' /data/frog.jpg')
3 my_show(plt.gca(), ml)
4 min_r, min_c = (min(ml.shape[0], frog.shape[0]),
5                 min(ml.shape[1], frog.shape[1])) # 取得可叠加区域
6
7 # blending of two images:
8 # by: img1 * wgt1 + img2 * wgt2 + wgt3
9 #      addWeights(img1, wgt1, img2, wgt2, wgt3)
10 dst = cv2.addWeighted( ml[:min_r, :min_c], 0.7,
11                        frog[:min_r, :min_c], 0.3, 0) #加权混合
12 my_show(plt.gca(), dst)
```



Histograms

```
1 apple = my_read_g(img_dir+'data/apple.png')
2 hist = cv2.calcHist([apple], [0], None, [256], [0,256]) # src imgs, color channels, mask
3                                                    # num bins, range
4 plt.hist(range(256), weights=hist, bins=256)
5 print(hist.shape)
```

(256, 1)



Histogram Equalization






```
1 apple = my_read_g(img_dir+'data/apple.png')
2 # reduce contrast (squash intensities)
3 # new min: 100, new max: 175
4 new = np.interp(apple, [apple.min(), apple.max()], [125, 175]).astype(np.uint8)
5 # equalize using CDF technique
6 equalized = cv2.equalizeHist(new)
7 fig, axes = plt.subplots(2,3,figsize=(12,6))
8 for idx, an_apple in enumerate([apple, new, equalized]):
9     # vmin/vmax set enforced min/max gray scale values ... without them
10    # 125 -> 0 ... 175 --> 255 and linearly interpolated
11    print("min: {} max: {}".format(an_apple.min(), an_apple.max()))
12    my_gshow(axes[0, idx], an_apple, vmin=0, vmax=255)
13    hist = cv2.calcHist([an_apple], [0], None, [256], [0, 256])
14    axes[1, idx].plot(hist)
```

min: 16 max: 241

min: 125 max: 175

min: 0 max: 255








```
[46]: 1 fig, axes = plt.subplots(2, 2, figsize=(12, 4))
      2 axes = axes.flat
      3
      4 # Otsu's thresholding
      5 next(axes).hist(blurred.flatten(), 256, log=True)
      6 # this seems weird: shouldn't it be black and white?!?
      7 thresh, th_otsu = cv2.threshold(blurred, 0, 255, cv2.THRESH_BINARY+cv2.THRESH_OTSU)
      8 print("Optimal Threshold is", thresh)
      9 my_show(next(axes), th_otsu, cmap='gray', interpolation=None) # force us to see what's there!
     10
     11 reblurred = cv2.GaussianBlur(blurred, (5, 5), 0) # neighborhood, variance?
     12 next(axes).hist(reblurred.flatten(), 256);
     13 thresh, th_otsu = cv2.threshold(reblurred, 0, 255, cv2.THRESH_BINARY+cv2.THRESH_OTSU)
     14 print("Optimal Threshold is", thresh)
     15 my_show(next(axes), th_otsu, cmap='gray') # also: interpolation=None)
```

Optimal Threshold is 117.0

Optimal Threshold is 126.0



Filters and Convolutions

```
1  # simple averaging filter without scaling parameter
2  mean_filter = np.ones((3,3))
3
4  # creating a gaussian filter
5  gk = cv2.getGaussianKernel(5,10)
6  gaussian = gk*gk.T
7
8  # different edge detecting filters
9  # laplacian
10 laplacian=np.array([[0, 1, 0],
11                    [1,-4, 1],
12                    [0, 1, 0]])
13
14 filters = [mean_filter, gaussian, laplacian, sobel_x, sobel_y, scharr]
15 filter_names = ['mean filter', 'gaussian', 'laplacian',
16                'sobel x', 'sobel y', 'scharr x']
17
18 fig, axes = plt.subplots(2,3,figsize=(8,5))
19 for name, filt, ax in zip(filter_names, filters, axes.flat):
20     # interesting variations:
21     # cmap='gray', 'jet', default; interpolation = None
22     my_show(ax, mag_fft(filt), cmap='jet')
23     ax.set_title(name)
```



作业课后布置

