

Pandas (续) C06



胡俊峰 2020/03/23

北京大学信息科学技术学院

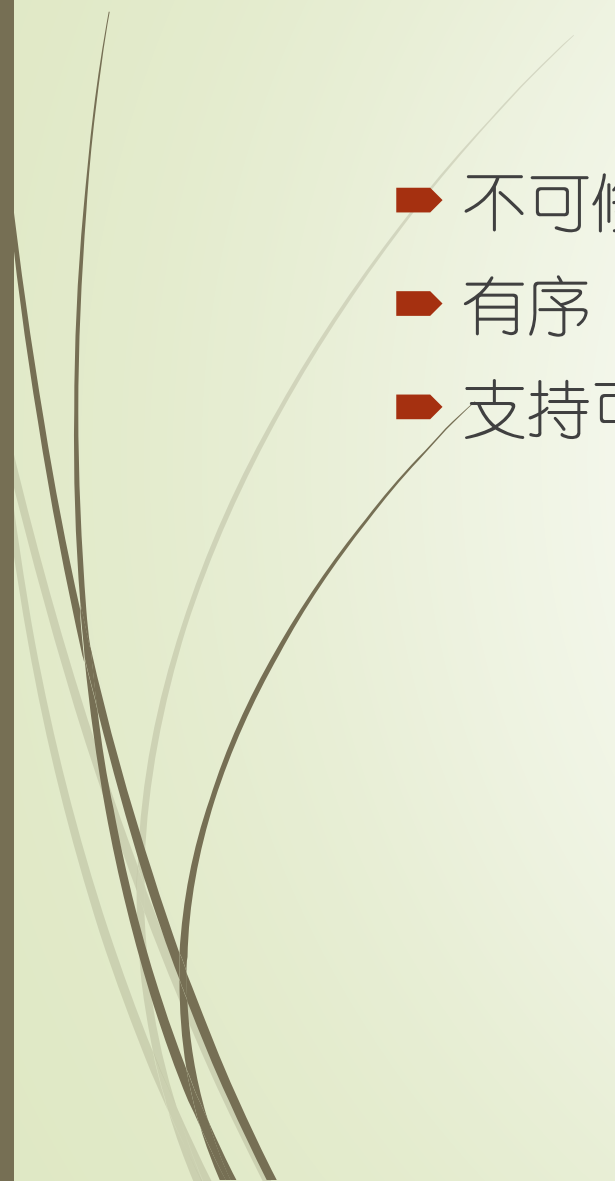


主要内容

- Pandas索引对象与高维索引
- 推荐算法与搜索排名
- Python交互界面与Tkinter软件包（郭一诺助教）



The Pandas Index Object

- 不可修改的数组
 - 有序
 - 支持可重复 key
- 

Index as ordered set (支持表关联计算的基础)

```
1 indA = pd.Index([1, 3, 5, 7, 9])
2 indB = pd.Index([2, 3, 5, 7, 11])
```

```
1 indA & indB # intersection
```

```
Int64Index([3, 5, 7], dtype='int64')
```

```
1 indA | indB # union
```

```
Int64Index([1, 2, 3, 5, 7, 9, 11], dtype='int64')
```

```
1 indA ^ indB # symmetric difference
```

```
Int64Index([1, 2, 9, 11], dtype='int64')
```

Indexers: loc, iloc, and ix

按位置索引

These slicing and indexing conventions can be a source of confusion. For example, if your `Series` has an explicit integer index, an indexing operation such as `data[1]` will use the explicit indices, while a slicing operation like `data[1:3]` will use the implicit Python-style index.

```
1 data = pd.Series(['a', 'b', 'c'], index=[1, 3, 5])
2 data
```

```
1    a
3    b
5    c
dtype: object
```

```
1 # explicit index when indexing
2 data[1]
```

```
'a'
```

```
1 # implicit index when slicing
2 data[1:3]
```

```
3    b
5    c
dtype: object
```

```
1 print(data.loc[1])
2 print(data.iloc[1])
3
```

```
a
b
```

Data Selection in DataFrame

```
1 area = pd.Series({'California': 423967, 'Texas': 695662,  
2                  'New York': 141297, 'Florida': 170312,  
3                  'Illinois': 149995})  
4 pop = pd.Series({'California': 38332521, 'Texas': 26448193,  
5                  'New York': 19651127, 'Florida': 19552860,  
6                  'Illinois': 12882135})  
7 data = pd.DataFrame({'area':area, 'pop':pop})  
8 data
```

	area	pop
California	423967	38332521
Florida	170312	19552860
Illinois	149995	12882135
New York	141297	19651127
Texas	695662	26448193

```
1 data['area']
```

```
California    423967
Florida       170312
Illinois      149995
New York      141297
Texas         695662
Name: area, dtype: int64
```

用法：字段名 类比 属性

Equivalently, we can use attribute-style access with column names that are strings:

```
1 data.area
```

```
California    423967
Florida       170312
Illinois      149995
New York      141297
Texas         695662
Name: area, dtype: int64
```

```
1 data['density'] = data['pop'] / data['area']  
2 data
```

	area	pop	density
California	423967	38332521	90.413926
Florida	170312	19552860	114.806121
Illinois	149995	12882135	85.883763
New York	141297	19651127	139.076746
Texas	695662	26448193	38.018740

行列互换

```
1 s = data.T
2 print(s)
3 s['California']
```

	California	Texas	New York	Florida	Illinois
area	4.239670e+05	6.956620e+05	1.412970e+05	1.703120e+05	1.499950e+05
pop	3.833252e+07	2.644819e+07	1.965113e+07	1.955286e+07	1.288214e+07
density	9.041393e+01	3.801874e+01	1.390767e+02	1.148061e+02	8.588376e+01

```
area      4.239670e+05
pop       3.833252e+07
density   9.041393e+01
```

```
Name: California. dtype: float64
```

筛选，赋值：

```
1 data.loc[data.density > 100, ['pop', 'density']]
```

	pop	density
Florida	19552860	114.806121
New York	19651127	139.076746

Any of these indexing conventions may also be used to set or modify values; this is done in the standard way that you might be accustomed to from working with NumPy:

```
1 data.iloc[0, 2] = 90  
2 data
```

	area	pop	density
California	423967	38332521	90.000000
Texas	695662	26448193	38.018740
New York	141297	19651127	139.076746
Florida	170312	19552860	114.806121
Illinois	149995	12882135	85.883763

Working with NumPy ufunc

```
1 df = pd.DataFrame(rng.randint(0, 10, (3, 4)),  
2                       columns=['A', 'B', 'C', 'D'])  
3 df
```

	A	B	C	D
0	9	2	6	7
1	4	3	7	7
2	2	5	4	1

```
1 np.sin(df * np.pi / 4)
```

采用Numpy的Ufunc —— 广播机制

	A	B	C	D
0	7.071068e-01	1.000000	-1.000000e+00	-0.707107
1	1.224647e-16	0.707107	-7.071068e-01	-0.707107
2	1.000000e+00	-0.707107	1.224647e-16	0.707107

Dataframe之间的运算自动进行索引对齐-补足 (out join)

Out[22]:

	A	B
0	2	4
1	18	6

► In [23]:

```
1 B = pd.DataFrame(rng.randint(0, 10, (3, 3)),  
2                       columns=list('BAC'))  
3 B
```

Out[23]:

	B	A	C
0	4	8	6
1	1	3	8
2	1	9	8

► In [24]:

```
1 A + B
```

Out[24]:

	A	B	C
0	10.0	8.0	NaN
1	21.0	7.0	NaN
2	NaN	NaN	NaN

The following table lists Python operators and their equivalent Pandas object methods:

Python Operator	Pandas Method(s)
<code>+</code>	<code>add()</code>
<code>-</code>	<code>sub()</code> , <code>subtract()</code>
<code>*</code>	<code>mul()</code> , <code>multiply()</code>
<code>/</code>	<code>truediv()</code> , <code>div()</code> , <code>divide()</code>
<code>//</code>	<code>floordiv()</code>
<code>%</code>	<code>mod()</code>
<code>**</code>	<code>pow()</code>

Frame 与 series 计算, 按行broadcasting

```
: 1 A = rng.randint(10, size=(3, 4))  
2 A
```

```
: array([[9, 4, 1, 3],  
        [6, 7, 2, 0],  
        [3, 1, 7, 3]])
```

```
: 1 df = pd.DataFrame(A, columns=list('QRST'))  
2 df - df.iloc[0]
```

	Q	R	S	T
0	0	0	0	0
1	-3	3	1	-3
2	-6	-3	6	0

```
1 df.subtract(df['R'], axis=0)
```

	Q	R	S	T
0	5	0	-3	-1
1	-1	0	-5	-7
2	2	0	6	2

运算过程中类型自适应转换

The following table lists the upcasting conventions in Pandas when NA values are introduced:

Typeclass	Conversion When Storing NAs	NA Sentinel Value
floating	No change	np. nan
object	No change	None or np. nan
integer	Cast to float64	np. nan
boolean	Cast to object	None or np. nan

Keep in mind that in Pandas, string data is always stored with an `object` dtype.

Detecting null values

Pandas data structures have two useful methods for detecting null data: `isnull()` and `notnull()`. Either one will return a Boolean mask over the data. For example:

```
data = pd.Series([1, np.nan, 'hello', None])
```

```
data.isnull()
```

```
0    False
1     True
2    False
3     True
dtype: bool
```

As mentioned in [Data Indexing and Selection](#), Boolean masks can be used directly as a `Series` or `DataFrame` index:

```
data[data.notnull()]
```

```
0      1
2  hello
dtype: object
```


We can fill NA entries with a single value, such as zero:

```
data.fillna(0)
```

```
a    1.0  
b    0.0  
c    2.0  
d    0.0  
e    3.0  
dtype: float64
```

We can specify a forward-fill to propagate the previous value forward:

```
# forward-fill  
data.fillna(method='ffill')
```

```
a    1.0  
b    1.0  
c    2.0  
d    2.0  
e    3.0  
dtype: float64
```

层次-组合 索引 (Hierarchical-Indexing)

```
1 index = [('California', 2000), ('California', 2010),
2         ('New York', 2000), ('New York', 2010),
3         ('Texas', 2000), ('Texas', 2010)]
4 populations = [33871648, 37253956,
5               18976457, 19378102,
6               20851820, 25145561]
7 pop = pd.Series(populations, index=index)
8 pop
```

```
(California, 2000)    33871648
(California, 2010)    37253956
(New York, 2000)      18976457
(New York, 2010)      19378102
(Texas, 2000)         20851820
(Texas, 2010)         25145561
dtype: int64
```

类似二维表切片

```
1 pop[:, 2010]

California    37253956
New York      19378102
Texas         25145561
dtype: int64
```

MultilIndex VS extra dimension

```
1 #unstack() method will quickly convert a multiply indexed Series
2 #into a conventionally indexed DataFrame:
3 pop_df = pop.unstack() ←
4 pop_df
```

	2000	2010
California	33871648	37253956
New York	18976457	19378102
Texas	20851820	25145561

```
1 #unstack() method will quickly convert a multiply indexed Series into a conventi
2 pop_df.stack()
```

California	2000	33871648
	2010	37253956
New York	2000	18976457
	2010	19378102
Texas	2000	20851820
	2010	25145561

dtype: int64

```

1 pop_df = pd.DataFrame({'total': pop,
2                          'under18': [9267089, 9284094,
3                                       4687374, 4318033,
4                                       5906301, 6879014]})
5 pop_df

```

		total	under18
California	2000	33871648	9267089
	2010	37253956	9284094
New York	2000	18976457	4687374
	2010	19378102	4318033
Texas	2000	20851820	5906301
	2010	25145561	6879014

```

1 f_u18 = pop_df['under18'] / pop_df['total']
2 f_u18.unstack()

```

	2000	2010
California	0.273594	0.249211
New York	0.247010	0.222831
Texas	0.283251	0.273568

Methods of MultiIndex Creation

The most straightforward way to construct a multiply indexed `Series` or `DataFrame` is to simply pass a list of two or more index arrays to the constructor. For example:

```
2]: 1 df = pd.DataFrame(np.random.rand(4, 2),  
2      index=[['a', 'a', 'b', 'b'], [1, 2, 1, 2]],  
3      columns=['data1', 'data2'])  
4 df
```

```
]:
```

		data1	data2
a	1	0.554233	0.356072
	2	0.925244	0.219474
b	1	0.441759	0.610054
	2	0.171495	0.886688

Similarly, if you pass a dictionary with appropriate tuples as keys, Pandas will automatically recognize this and use a `MultiIndex` by default:

```
1 data = {('California', 2000): 33871648,  
2         ('California', 2010): 37253956,  
3         ('Texas', 2000): 20851820,  
4         ('Texas', 2010): 25145561,  
5         ('New York', 2000): 18976457,  
6         ('New York', 2010): 19378102}  
7 pd.Series(data)
```

California	2000	33871648
	2010	37253956
New York	2000	18976457
	2010	19378102
Texas	2000	20851820
	2010	25145561

dtype: int64

MultiIndex constructor

```
1 pd.MultiIndex.from_arrays(['a', 'a', 'b', 'b'], [1, 2, 1, 2])
```

```
MultiIndex(levels=(['a', 'b'], [1, 2]),  
            labels=[[0, 0, 1, 1], [0, 1, 0, 1]])
```

You can construct it from a list of tuples giving the multiple index values of each point:


```
1 pd.MultiIndex.from_tuples([('a', 1), ('a', 2), ('b', 1), ('b', 2)])
```

```
MultiIndex(levels=(['a', 'b'], [1, 2]),  
            labels=[[0, 0, 1, 1], [0, 1, 0, 1]])
```

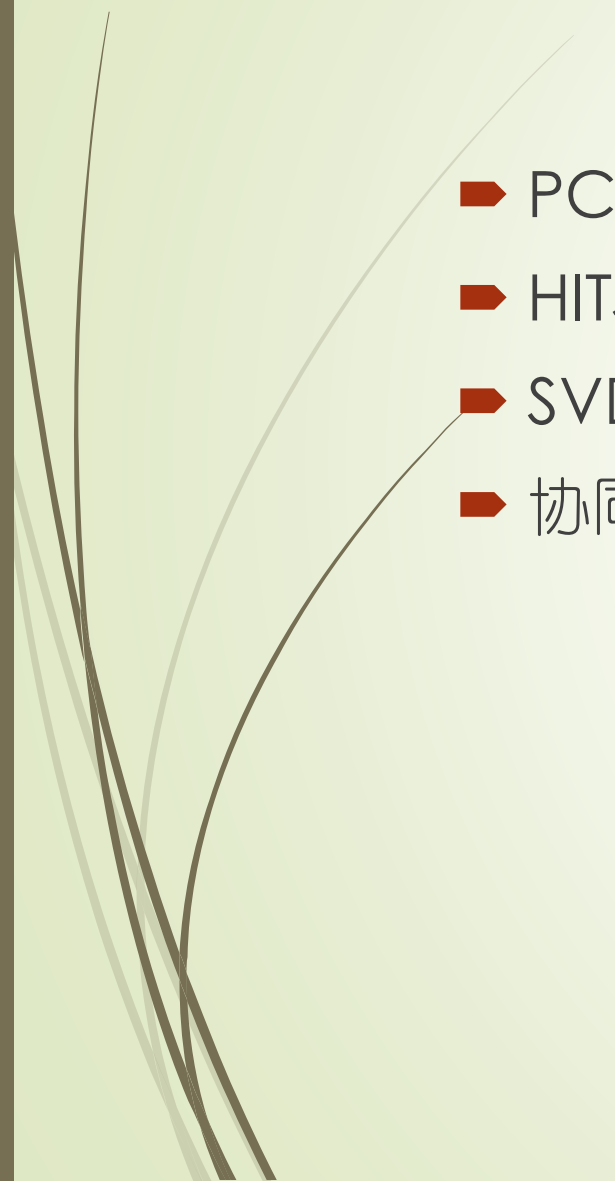
You can even construct it from a Cartesian product of single indices:

```
1 pd.MultiIndex.from_product(['a', 'b'], [1, 2])
```

```
MultiIndex(levels=(['a', 'b'], [1, 2]),  
            labels=[[0, 0, 1, 1], [0, 1, 0, 1]])
```

推荐算法与搜索排名

- PCA降维与特征空间
 - HITS算法与搜索排名
 - SVD与隐含语义挖掘
 - 协同过滤算法
- 

HITS算法(Hyperlink-induced Topic Search)

网络分析基本
方法及其应用

张涵

网络分析基本
思想

静态结构分
析：基础统计

基础统计

社会群体及社会角色

弱关系及桥

点的同质

性(homophily): 强关
系的影响

对强弱关系的反思

对结构的深入
研究：建立模
型

网络平衡

度数分布的深入研

究：模型的实例

网络的链接分析及预
测

数据收集及处

- 首先，通过关键词在文档中是否出现，得到相关网页集合及其链接关系。
- 每个网页有两个值： a (权威值，入度高)及 h (hub, 中枢值，出度高)
- 通过反复迭代计算，得出得分最高的网页。



定义n个网页之间链接关系的邻接矩阵为M, 即 M_{ij} 为 $1 \iff$ 从网页i到j有链接, 则网页i的中枢值为:

$$a_i \leftarrow M_{1i}h_1 + M_{2i}h_2 + \dots M_{ni}h_n \quad (3)$$

$$h_i \leftarrow M_{i1}a_1 + M_{i2}a_2 + \dots M_{in}a_n \quad (4)$$

$h^k = (h_1, h_2, \dots, h_n)^T$, $a^k = (a_1, a_2, \dots, a_n)^T$ 为运行了k次的时候, n个网页的中枢, 权威向量, 则转换规则为: (先更新 a^k)

$$a^k = M^T h^{k-1} \quad (5)$$

$$h^k = M a^k \quad (6)$$

h^0, a^0 的元素都为1 迭代可得: $a^1 = M^T h^0, h^1 = M M^T h^0, \dots$

$$a^k = (M^T M)^{k-1} h^0 \quad (7)$$

$$h^k = (M M^T)^k h^0 \quad (8)$$

定理: $n \times n$ 的实对称矩阵有 n 个特征值, 且不同特征值的特征向量彼此正交 (即特征向量构成线性空间的一组基底)

设 $(M^T M)^{k-1}$ 的特征值为 c_1, c_2, \dots, c_n , 且 $c_1 > c_2 > \dots > c_n$, 对应的特征向量为 z_1, z_2, \dots, z_n , 而 h^0 在基底下的表示为

$$h_0 = q_1 z_1 + q_2 z_2 + \dots + q_n z_n \quad (9)$$

则

$$h^k = (MM^T)^k h^0 \quad (10)$$

$$= (MM^T)^k (q_1 z_1 + q_2 z_2 + \dots + q_n z_n) \quad (11)$$

$$= q_1 (MM^T)^k z_1 + q_2 (MM^T)^k z_2 + \dots + q_n (MM^T)^k z_n \quad (12)$$

$$= q_1 c_1^k z_1 + q_2 c_2^k z_2 + \dots + q_n c_n^k z_n \quad (13)$$

如果要收敛，需要对每项正规化。则

$$h^k = \frac{(MM^T)^k h^0}{\|(MM^T)^k h^0\|} \quad (14)$$

$$= \frac{q_1 c_1^k z_1 + q_2 c_2^k z_2 + \dots q_n c_n^k z_n}{\|q_1 c_1^k z_1 + q_2 c_2^k z_2 + \dots q_n c_n^k z_n\|} \quad (15)$$

$$= \frac{q_1 z_1 + q_2 (\frac{c_2}{c_1})^k z_2 + \dots q_n (\frac{c_n}{c_1})^k z_n}{\|q_1 z_1 + q_2 (\frac{c_2}{c_1})^k z_2 + \dots q_n (\frac{c_n}{c_1})^k z_n\|} \quad (16)$$

$$= \frac{q_1 z_1}{\|q_1 z_1\|} \quad (17)$$

故 h_k 收敛，同理可得 a_k 收敛

奇异值分解

➤ $N \times M$ 矩阵 X ，把每一行当成一个点。设 $r = \text{rank}(X)$ 。

➤ $\vec{v}_1 = \operatorname{argmax}_{|\vec{v}|=1} |X\vec{v}|$

➤ $\vec{v}_2 = \operatorname{argmax}_{|\vec{v}|=1, \vec{v} \perp \vec{v}_1} |X\vec{v}|$

➤

➤ $\vec{v}_r = \operatorname{argmax}_{|\vec{v}|=1, \vec{v} \perp \vec{v}_1, \vec{v} \perp \vec{v}_2, \dots, \vec{v} \perp \vec{v}_{r-1}} |X\vec{v}|$

➤ 令 $\sigma_i = |X\vec{v}_i|$, $\vec{u}_i = \frac{1}{\sigma_i} X\vec{v}_i$, $1 \leq i \leq r$

➤ 则 $X = \sum_{i=1}^r \sigma_i \vec{u}_i \vec{v}_i^T = \begin{bmatrix} | & & | \\ \vec{u}_1 & \dots & \vec{u}_r \\ | & & | \end{bmatrix} \begin{bmatrix} \sigma_1 & & \\ & \ddots & \\ & & \sigma_r \end{bmatrix} \begin{bmatrix} -\vec{v}_1^T & - \\ \dots & \\ -\vec{v}_r^T \end{bmatrix} = UDV^T$

奇异值分解

➤ $N \times M$ 矩阵 X ，把每一行当成一个点。设 $r = \text{rank}(X)$ 。


$$\text{➤ } X = \sum_{i=1}^r \sigma_i \vec{u}_i \vec{v}_i^T = \begin{bmatrix} | & & | \\ \vec{u}_1 & \dots & \vec{u}_r \\ | & & | \end{bmatrix} \begin{bmatrix} \sigma_1 & & \\ & \ddots & \\ & & \sigma_r \end{bmatrix} \begin{bmatrix} -\vec{v}_1^T & - \\ \dots & \\ -\vec{v}_r^T \end{bmatrix} = UDV^T$$

➤ U, D, V 分别为 $N \times r, r \times r, M \times r$ 矩阵

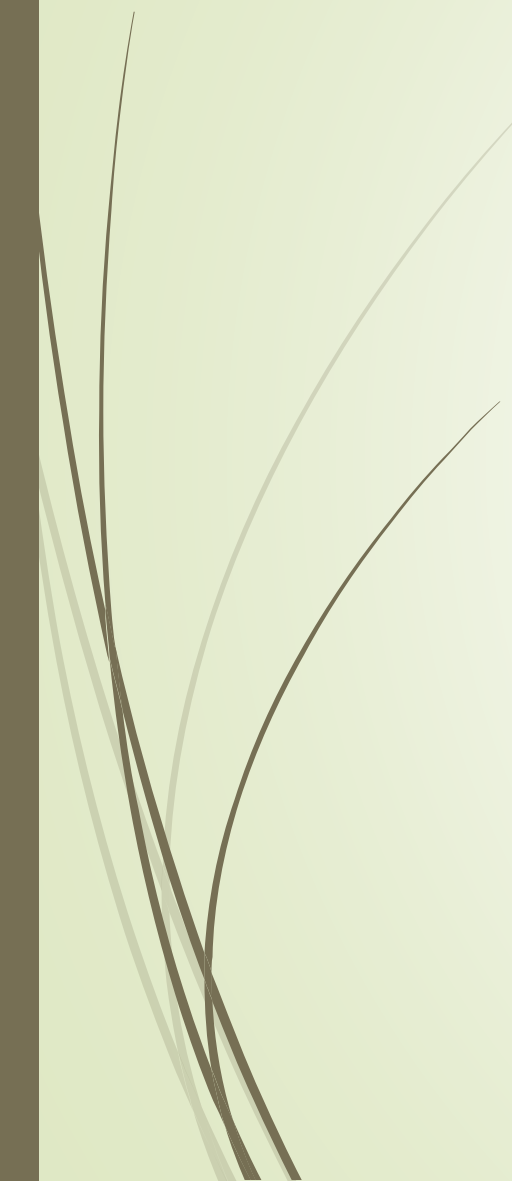
➤ $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_r$ 称为奇异值


奇异值分解与主成分分析

- $X = UDV^T$
- 考虑 X 的变换 $Y = XV = UD$ (维度重组)
- 则 $C_Y = \frac{1}{N}Y^TY = \frac{1}{N} = \frac{1}{N}D^T U^T U D = \frac{1}{N}D^2$ 是个对角矩阵
- 各个维度之间不相关
- 对角元的大小——这一维自身的方差
- 按方差大小排列, 得到第 $1, 2, \dots, r$ 个主成分
- 方差小的几个主成分可以认为是噪声, 将其丢弃——降维



降维的意义

- 减少数据冗余度
 - 一定程度上减少了噪声
- 



一个直接的应用

- 隐含语义挖掘 (LSI 或称 LSA)

T-SNE T分布随机近邻嵌入

1. PCA和LDA都是线性的，t-sne是一种非线性的降维算法，非常适用于将高维数据降维到2维和3维，进行可视化。
2. SNE构建一个高维对象之间的概率分布，使得相似的对象有更高的概率被选择，而不相似的对象有较低的概率被选择。
3. SNE在低维空间里在构建这些点的概率分布，使得这两个概率分布之间尽可能的相似。

```

from sklearn.manifold import TSNE
from matplotlib import patheffects as PathEffects
tsne = TSNE(n_components=2)
tsne_results = tsne.fit_transform(X_std)

def scatter(x, colors):

    # We create a scatter plot.
    f = plt.figure(figsize=(8, 8))
    ax = plt.subplot(aspect='equal')
    sc = ax.scatter(x[:,0], x[:,1], lw=0, s=40,
                    c = Target.values, cmap='jet',)


    plt.xlim(-25, 25)
    plt.ylim(-25, 25)
    ax.axis('off')
    ax.axis('tight')


    # We add the labels for each digit.
    txts = []
    for i in range(10):
        # Position of each label.
        xtext, ytext = np.median(x[colors == i, :], axis=0)
        txt = ax.text(xtext, ytext, str(i), fontsize=24)
        txt.set_path_effects([
            PathEffects.Stroke(linewidth=5, foreground="w"),
            PathEffects.Normal()])
        txts.append(txt)

    return f, ax, sc, txts

scatter(tsne_results, Target.values)

```





作业群里布置：

